

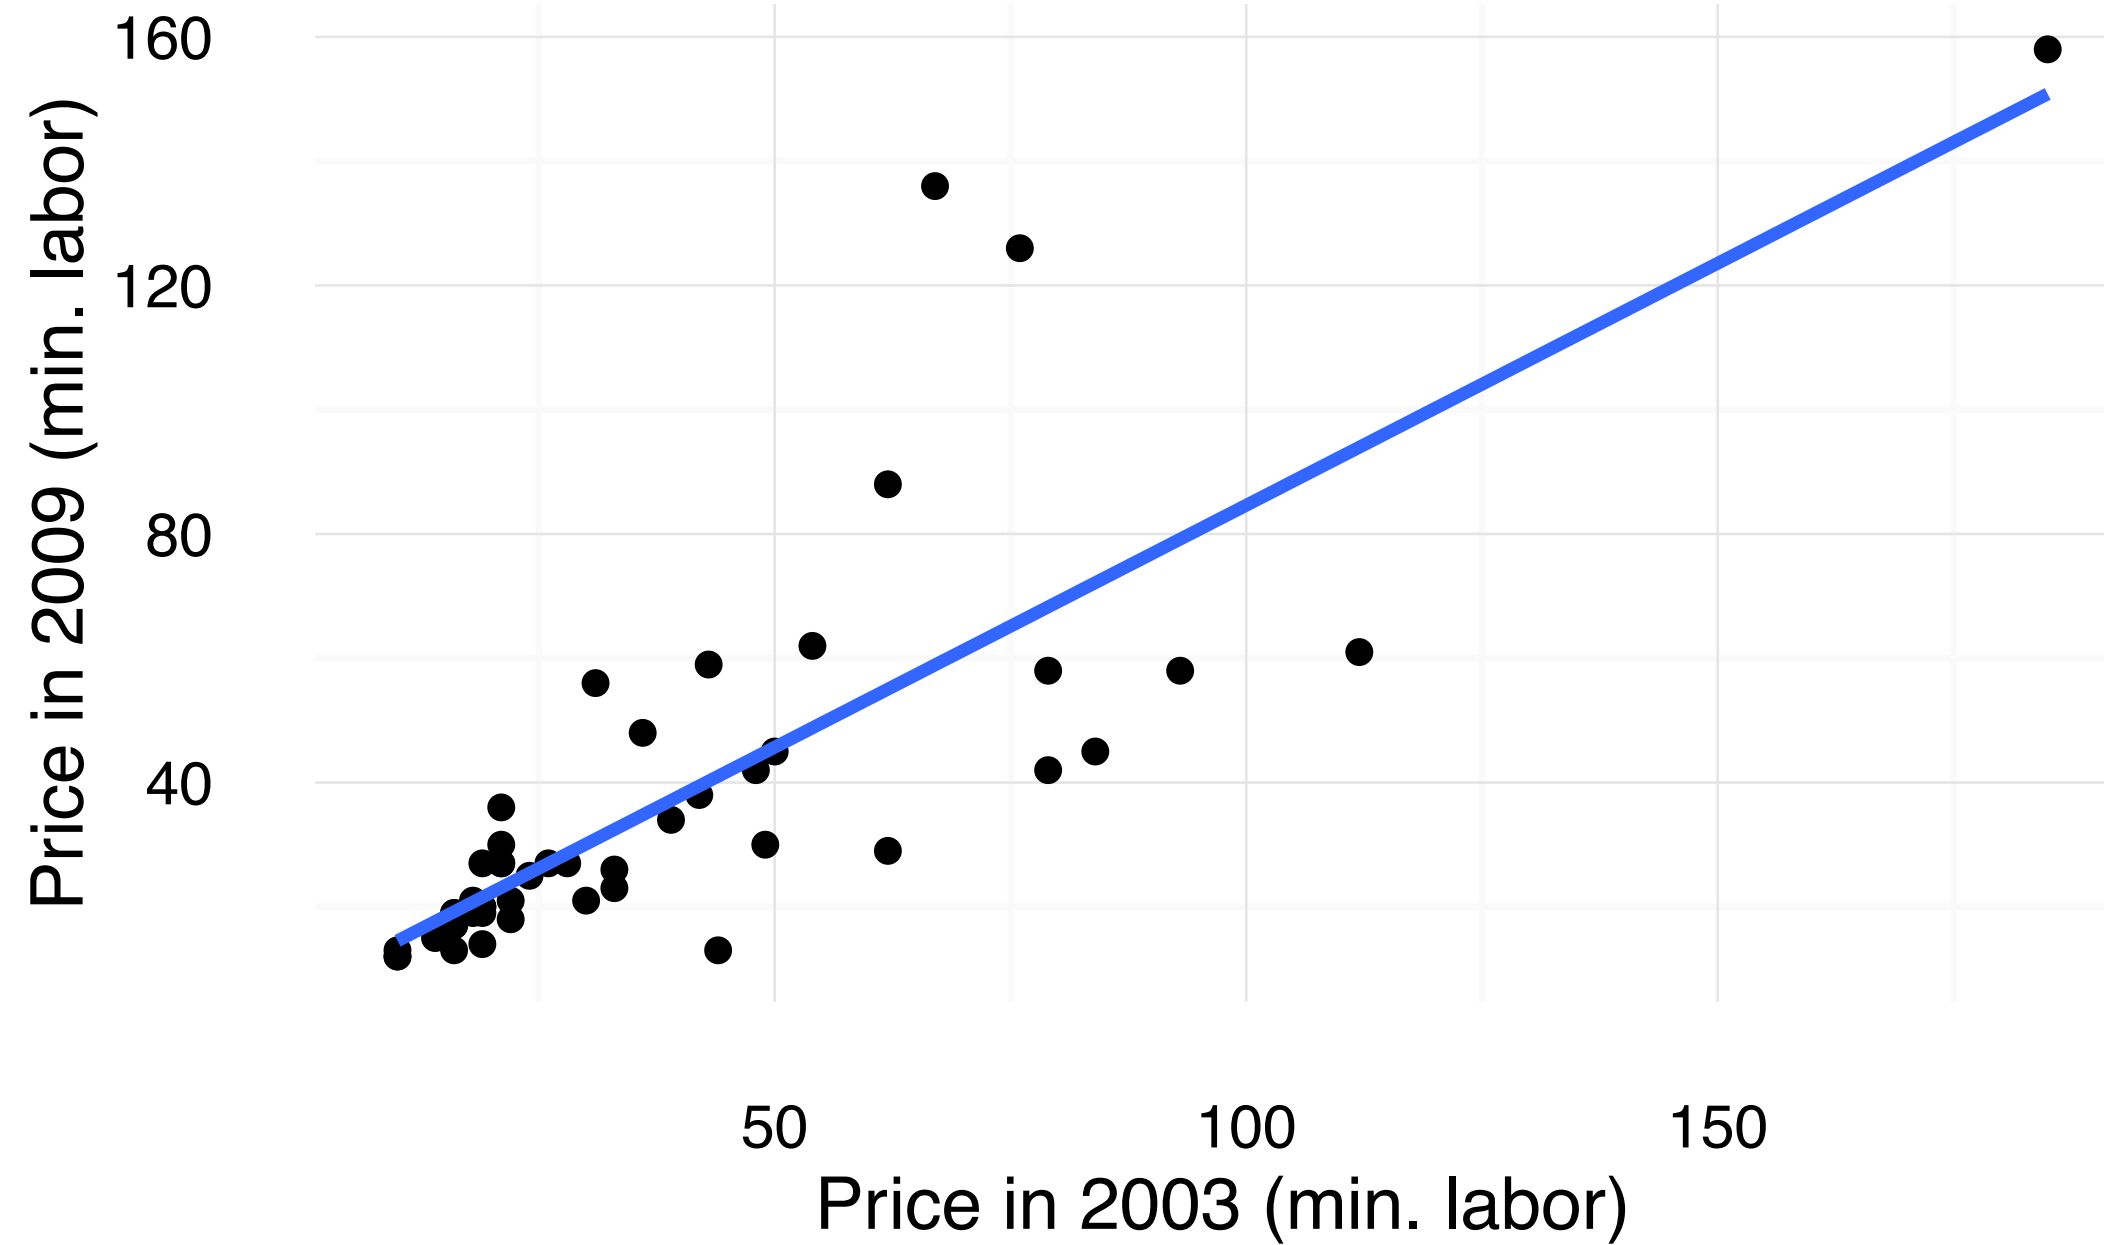
Assessing the Model

Union Bank of Switzerland (UBS) reports

Produces regular reports on prices & earnings in major cities throughout the world, in minutes of labor required for "typical" worker to purchase the commodity:

- prices of basic commodities (1 kg rice, 1 kg bread)**
- price of a Big Mac at McDonald's**

Data from 2003 (before recession) and 2009 (after recession) reports



```
cor(bigmac2009 ~ bigmac2003, data = UBSprices)
[1] 0.804481
```

```
bigmac.lm <- lm(bigmac2009 ~ bigmac2003, data = UBSprices)
summary(bigmac.lm)
```

Call:

```
lm(formula = bigmac2009 ~ bigmac2003, data = UBSprices)
```

Residuals:

Min	1Q	Median	3Q	Max
-32.968	-5.258	-2.159	0.187	77.081

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.73612	3.84985	1.750	0.0861 .
bigmac2003	0.77886	0.07975	9.767	2.33e-13 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18.35 on 52 degrees of freedom

Multiple R-squared: 0.6472, Adjusted R-squared: 0.6404

F-statistic: 95.39 on 1 and 52 DF, p-value: 2.334e-13

Bootstrapping **regression** models

- 1. Number cases in data set from 1 to n .**
- 2. Take a random sample *with replacement* of size n from these numbers.**
- 3. Create a new data set by pulling the rows (cases) from the original data set that were selected in the random sample.**
- 4. Fit the regression model to this new data set and save the values of the estimated coefficients or other summary statistics.**
- 5. Repeat steps 2-4 R times.**

Bootstrap confidence intervals

95% plug-in intervals

Percentile intervals

Plug-in intervals in R

```
bigmac.boot <- Boot(bigmac.lm, R = 999)
```

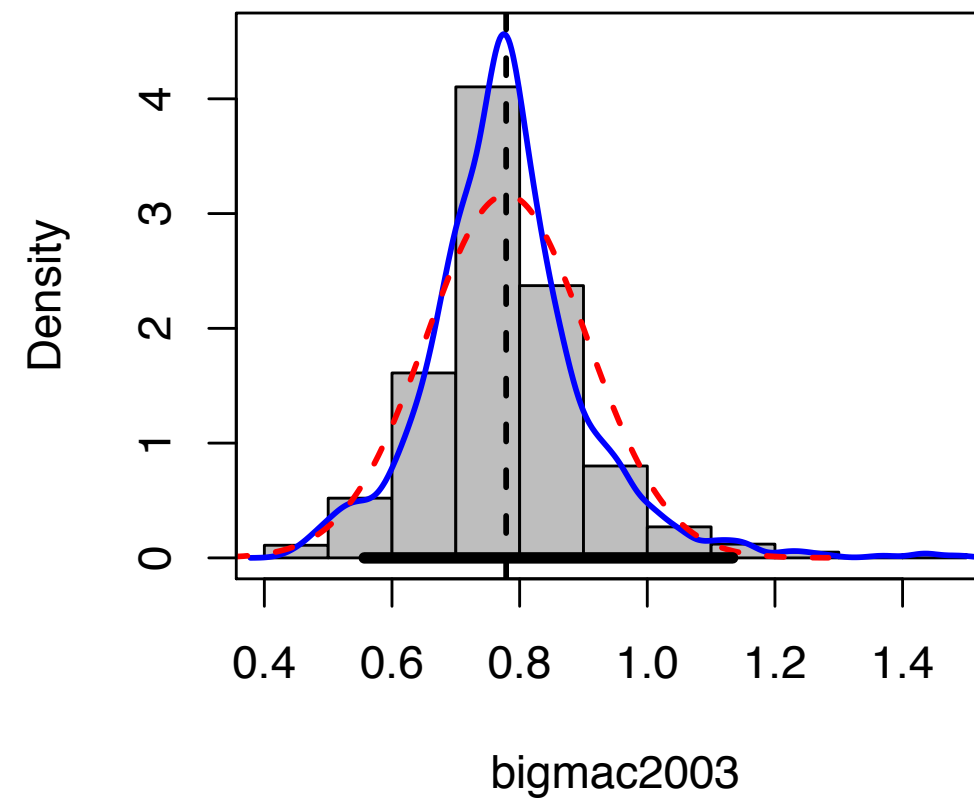

```
summary(bigmac.boot)
```

	R	original	bootBias	bootSE	bootMed
(Intercept)	999	6.73612	-0.01700515	2.84323	6.69652
bigmac2003	999	0.77886	0.00049195	0.12613	0.77432

Interpreting confidence intervals for the slope

Are the bootstrap distributions bell-shaped?

```
hist(bigmac.boot, col = "gray")
```



Percentile intervals in R

```
confint(bigmac.boot, type = "perc", level = .95)
```

```
Bootstrap quantiles, type = percent
```

	2.5 %	97.5 %
(Intercept)	0.2406487	12.243993
bigmac2003	0.5309925	1.058041

Is the model
any good?

Coefficient of determination: R^2

Interpretation: Proportion of the variability in y explained by the linear model.

Calculation: For single variable regression, $R^2 = r^2$

Intuition: A better model explains more of the variability in y

Pitfall: R^2 does not talk about predictive ability of the model

```
summary(bigmac.lm)
```

```
Call:
```

```
lm(formula = bigmac2009 ~ bigmac2003, data = UBSprices)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-32.968	-5.258	-2.159	0.187	77.081

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.73612	3.84985	1.750	0.0861 .
bigmac2003	0.77886	0.07975	9.767	2.33e-13 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 18.35 on 52 degrees of freedom
```

```
Multiple R-squared:  0.6472,    Adjusted R-squared:  0.6404
```

```
F-statistic: 95.39 on 1 and 52 DF,  p-value: 2.334e-13
```

Model assumptions


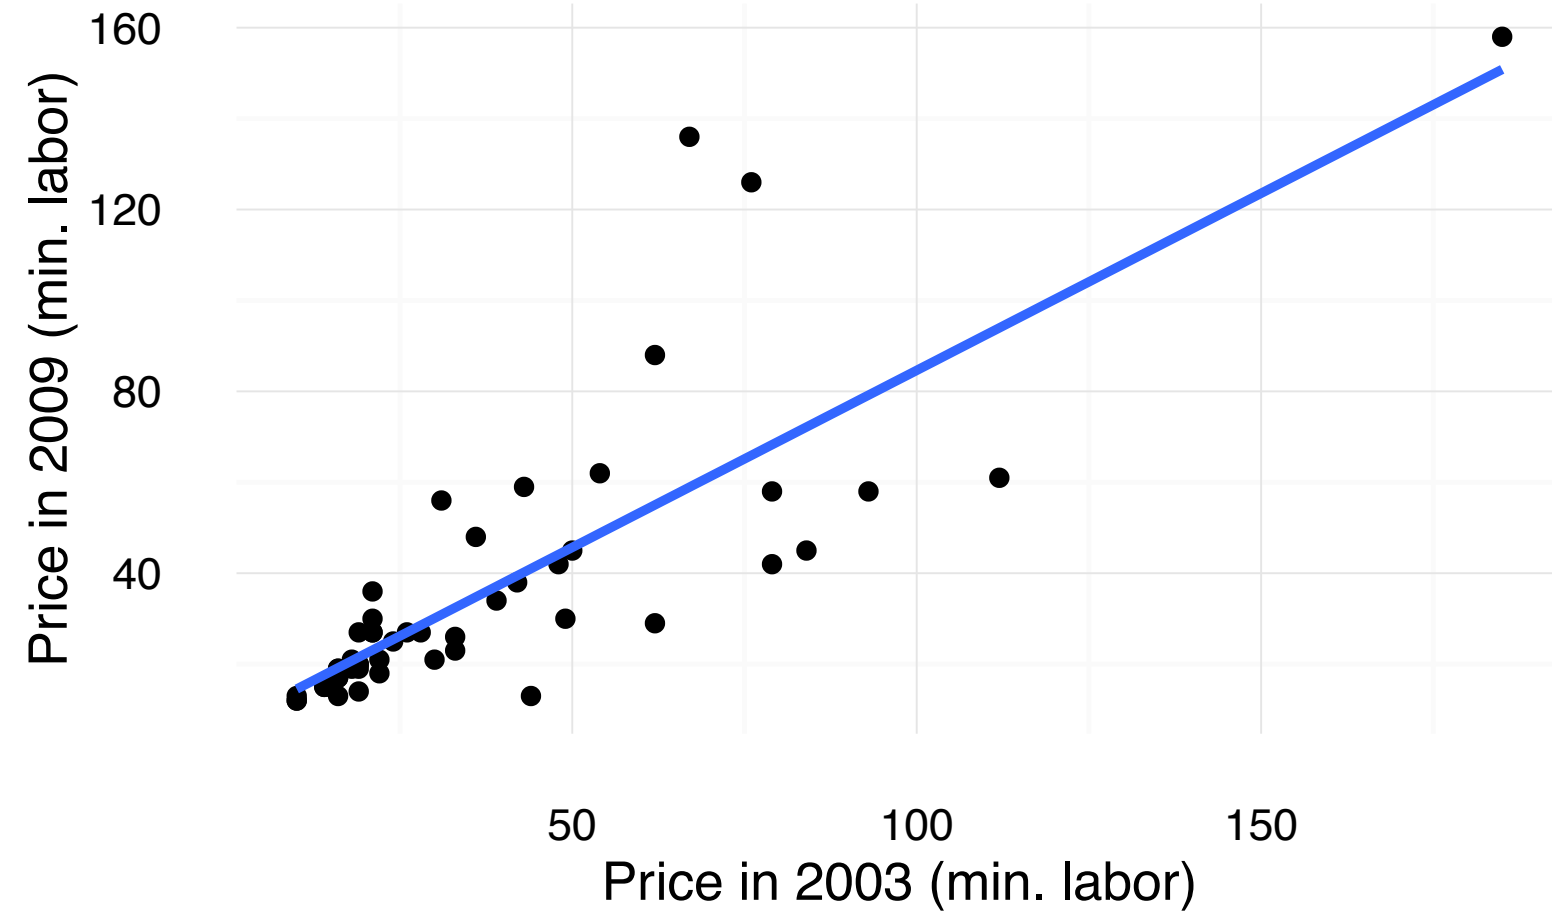
Linearity

Constant variability around the line

Bell-shaped residuals



Residuals

$$e_i = y_i - \hat{y}_i$$



Residuals

$$e_i = y_i - \hat{y}_i$$



Calibrate your intuition


https://gallery.shinyapps.io/slr_diag/

Outliers

Leverage point

Away from the cloud of points horizontally


Does not necessarily change the slope




Influential point

Changes the slope (most likely also has high leverage)

Run the regression with and without that point to determine



Outlier



An unusual point without these special characteristics (this one likely affects the intercept only)

If clusters (groups of points) are apparent in the data, it might be worthwhile to model the groups separately